

# RAG Document Ingestion & Retrieval — File & Folder Structure Note

## 1. Purpose

Define a **clean, scalable repository structure** that directly reflects system components and responsibilities.

This structure is designed to:

- Prevent tight coupling
- Make ownership clear
- Support future scaling (services / pipelines)
- Be understandable to mentors, PMs, and new engineers

No framework or language is assumed.

---

## 2. Core Principle (Industry Rule)

**Folders represent responsibilities, not technologies.**

If a folder name contains a tool name, the design is already leaking.

---

## 3. High-Level Repository Layout

```
rag-system/
|
|   docs/
|   |
|   |   design-note.md
|   |   implementation-approach.md
|   |   tools-component-mapping.md
|   |   file-structure.md
|
|   config/
|   |
|   |   rules/
|   |   |
|   |   |   parsing-rules.yaml
|   |   |   promotion-rules.yaml
|   |   |   trust-thresholds.yaml
|
|   |   environments/
|   |   |
|   |   |   dev.yaml
|   |   |   prod.yaml
|
|   core/
```

```
|   ├── registry/
|   ├── probe/
|   ├── decision_engine/
|   ├── parsing/
|   ├── quality/
|   ├── promotion/
|   └── orchestration/
|
|   ├── storage/
|   │   ├── sql/
|   │   ├── object_store/
|   │   ├── vector/
|   │   └── graph/
|
|   ├── retrieval/
|   │   ├── query_router/
|   │   ├── vector_search/
|   │   ├── graph_reasoning/
|   │   └── response_builder/
|
|   ├── workflows/
|   │   ├── ingestion_flow/
|   │   ├── parsing_flow/
|   │   └── promotion_flow/
|
|   ├── shared/
|   │   ├── models/
|   │   ├── schemas/
|   │   ├── utils/
|   │   └── logging/
|
|   ├── tests/
|   │   ├── unit/
|   │   ├── integration/
|   │   └── replay/
|
|   └── README.md
```

---

## 4. Folder Responsibilities (What Goes Where)

### 4.1 /docs

Design and decision artifacts.

**Why:** Architecture lives longer than code.

---

#### 4.2 /config

All **rules and thresholds** live here — not hardcoded.

- Parsing strategies
- Promotion thresholds
- Trust policies

**Why:** Allows tuning without redeploying code.

---

#### 4.3 /core

Contains **business logic**, independent of storage or tools.

Subfolders: - `registry/` → document lifecycle - `probe/` → structure inspection - `decision_engine/`  
→ logical parameters - `parsing/` → deterministic + ML routing - `quality/` → confidence & trust  
scoring - `promotion/` → routing to vector / graph - `orchestration/` → workflow coordination

**Rule:** No external system calls directly from here.

---

#### 4.4 /storage

All persistence logic.

- `sql/` → authoritative store
- `object_store/` → PDFs, images
- `vector/` → embeddings
- `graph/` → relationships

**Rule:** Storage layers are dumb and replaceable.

---

#### 4.5 /retrieval

User-facing query resolution.

- `query_router/` → decides vector vs graph
- `vector_search/` → semantic retrieval
- `graph_reasoning/` → relationship queries
- `response_builder/` → grounded responses

**Guarantee:** Responses must reference SQL-backed data.

---

#### 4.6 /workflows

Defines **end-to-end flows**.

- Ingestion
- Parsing
- Promotion

**Why:** Keeps orchestration separate from logic.

---

#### 4.7 /shared

Common definitions.

- Data models
- Schemas
- Utilities
- Logging

**Rule:** No business decisions here.

---

#### 4.8 /tests

Testing mirrors architecture.

- Unit tests per component
  - Integration tests per pipeline
  - Replay tests for reprocessing documents
- 

### 5. Why This Structure Works

- Mirrors system architecture exactly
  - Prevents cross-layer leakage
  - Encourages clean abstractions
  - Makes scaling to microservices trivial
  - Easy for reviewers to understand
-

## 6. Migration & Scaling Path

- Start as a modular monolith
  - Extract `/core` subfolders into services later
  - Storage adapters remain unchanged
  - Workflows evolve independently
- 

## 7. One-Sentence Summary

**"This repository structure enforces trust, separation of concerns, and long-term maintainability by construction."**

---

*End of file & folder structure note.*